

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
%matplotlib inline
import seaborn as sns
sns.set()
sns.set_style("darkgrid")
import warnings
warnings.filterwarnings("ignore")
import math
from sklearn.metrics import accuracy_score
import statsmodels.api as sm
import numpy as np
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score, GridSearchCV, KFold
from sklearn import preprocessing
```

Question 1

The “Variable Description.xlsx” spreadsheet contains a list of variables that we’ll use for our analyses. Note that this is not a full list of all the variables in the dataset, although it’s close (we ignoring a few perfectly co-linear predictors). Filter the full set of variables in the dataset down to the the Opportunity Insights and PM COVID variables listed in the spreadsheet along with county, state, and deathspc.

```
In [2]: vd_data = pd.read_excel("PPHA_30545_MP03-Variable_Description.xlsx")

vd_data['Variable'].unique

vd_variables = vd_data['Variable'].tolist()

print(vd_variables)
```

```
['casespc', 'deathspc', 'intersects_msa', 'cur_smoke_q1', 'cur_smoke_q2', 'cur_smoke_q3', 'cur_smoke_q4', 'bmi_obese_q1', 'bmi_obese_q2', 'bmi_obese_q3', 'bmi_obese_q4', 'exercise_any_q1', 'exercise_any_q2', 'exercise_any_q3', 'exercise_any_q4', 'brfss_mia', 'puninsured2010', 'reimb_penroll_adj10', 'mort_30day_hosp_z', 'adjmortmeas_amiall130day', 'adjmortmeas_chfall130day', 'med_prev_qual_z', 'primcarevis_10', 'diab_hemotest_10', 'diab_eyeexam_10', 'diab_lipids_10', 'mammogram_10', 'cs00_seg_inc', 'cs00_seg_inc_pov25', 'cs00_seg_inc_aff75', 'cs_race_theil_2000', 'gini99', 'poor_share', 'inc_share_1perc', 'frac_middleclass', 'scap_ski90pcm', 'rel_tot', 'cs_frac_black', 'cs_frac_hisp', 'unemp_rate', 'cs_labforce', 'cs_elf_ind_man', 'cs_born_foreign', 'mig_inflow', 'mig_out']
```

```
flow', 'pop_density', 'frac_traveltime_lt15', 'hhinc00', 'median_house_value', 'ccd_exp_tot', 'score_r', 'cs_fam_wkidsinglemom',
'subcty_exp_pc', 'taxrate', 'tax_st_diff_top20', 'pm25', 'pm25_mia', 'summer_tmmx', 'summer_rmax', 'winter_tmmx', 'winter_rmax',
'bmcruerate']
```

```
In [3]: covid_data = pd.read_csv('Data-Covid002.csv', encoding='ISO-8859-1')
covid_data.columns
covid_data.set_index(['county'], inplace = True)
vd_variables.append('state')
covid_data1 = covid_data[vd_variables[1:]]
```

```
In [4]: covid_data1
```

```
Out[4]:
```

	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	bmi_obese
county										
Autauga	16.548864	1	0.333333	0.238095	0.208333	0.133333	0.375000	0.238095	0.260870	0.133
Baldwin	8.959118	1	0.268097	0.233503	0.167464	0.176991	0.298050	0.262467	0.193237	0.135
Barbour	6.609756	0	0.228571	0.250000	0.181818	0.111111	0.294118	0.571429	0.545455	0.277
Bibb	6.038192	1	0.244444	0.280000	0.181818	0.150000	0.466667	0.375000	0.190476	0.100
Blount	1.503713	1	0.304348	0.260870	0.352941	0.166667	0.347826	0.318182	0.529412	0.235
...
Sweetwater	0.000000	0	0.349304	0.302658	0.227799	0.178484	0.247209	0.226943	0.254593	0.257
Teton	10.535728	0	0.263415	0.182018	0.092105	0.053913	0.127451	0.099338	0.113333	0.072
Uinta	0.000000	0	0.345538	0.237718	0.197943	0.130682	0.283019	0.247154	0.224274	0.235
Washakie	14.836295	0	0.260163	0.195652	0.083333	0.076923	0.283186	0.139706	0.265306	0.194
Weston	0.000000	0	0.318841	0.241935	0.264706	0.266667	0.181818	0.290323	0.218750	0.285

3107 rows × 62 columns



Question 2

Compute descriptive (summary) statistics for the subset of Opportunity Insights and PM COVID variables you filtered in previous question.

In [5]:

```
print(covid_data1.describe())
```

	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	\
count	3107.000000	3107.000000	3107.000000	3107.000000	3107.000000	
mean	23.790131	0.596717	0.212659	0.171048	0.134467	
std	67.852145	0.490636	0.149348	0.128130	0.132181	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	3.802303	1.000000	0.250000	0.198718	0.142857	
75%	21.461759	1.000000	0.310931	0.250000	0.200000	
max	2279.610600	1.000000	1.000000	1.000000	1.000000	

	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	bmi_obese_q4	\
count	3107.000000	3107.000000	3107.000000	3107.000000	3107.000000	
mean	0.098316	0.239166	0.214580	0.209621	0.186739	
std	0.110110	0.165928	0.153237	0.175849	0.167227	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.080128	0.000000	0.000000	0.000000	
50%	0.096535	0.272076	0.241590	0.223124	0.194118	
75%	0.148719	0.335532	0.304348	0.297220	0.266667	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	...	subcty_exp_pc	taxrate	tax_st_diff_top20	pm25	\
count	...	3107.000000	3107.000000	3106.000000	3107.000000	
mean	...	2119.407531	0.023089	0.775634	8.371871	
std	...	999.833466	0.013848	1.470989	2.565927	
min	...	0.000000	0.000000	0.000000	0.000000	
25%	...	1510.192750	0.014993	0.000000	6.309710	
50%	...	1935.919400	0.020339	0.000000	8.784647	
75%	...	2505.411100	0.027164	1.000000	10.483764	
max	...	20541.918000	0.209907	7.220000	15.786018	

	pm25_mia	summer_tmmx	summer_rmax	winter_tmmx	winter_rmax	\
count	3107.000000	3107.000000	3107.000000	3107.000000	3107.000000	
mean	0.003540	303.126997	88.970517	280.404875	87.469432	
std	0.059405	3.173950	9.689271	6.597855	4.811207	
min	0.000000	290.455540	31.643282	264.693820	58.159798	
25%	0.000000	300.848035	88.052494	275.113020	85.093342	
50%	0.000000	303.290440	91.320313	280.154690	88.028793	
75%	0.000000	305.817430	94.812389	285.543750	90.747704	

```
max      1.000000   313.872680   99.778748   298.340360   97.672874
```

```
bmcruderate
count    3107.00000
mean     1029.15597
std       248.38181
min       189.30000
25%       864.29999
50%      1036.30000
75%      1194.10000
max       1978.60000
```

```
[8 rows x 61 columns]
```

Question 3

Note that some variables have missing values. This causes problems when estimating the models. Normally we'd impute missing values by replacing them with their mean or median value, but to keep things simple, given the size of our data, you should drop all observations (rows) with missing values.

```
In [6]: covid_data1.dropna(inplace=True)
```

```
In [7]: covid_data1
```

```
Out[7]:
```

	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	bmi_obese
county										
Autauga	16.548864	1	0.333333	0.238095	0.208333	0.133333	0.375000	0.238095	0.260870	0.133
Baldwin	8.959118	1	0.268097	0.233503	0.167464	0.176991	0.298050	0.262467	0.193237	0.135
Barbour	6.609756	0	0.228571	0.250000	0.181818	0.111111	0.294118	0.571429	0.545455	0.277
Bibb	6.038192	1	0.244444	0.280000	0.181818	0.150000	0.466667	0.375000	0.190476	0.100
Blount	1.503713	1	0.304348	0.260870	0.352941	0.166667	0.347826	0.318182	0.529412	0.235
...
Sweetwater	0.000000	0	0.349304	0.302658	0.227799	0.178484	0.247209	0.226943	0.254593	0.257

	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	bmi_obese
county										
Teton	10.535728	0	0.263415	0.182018	0.092105	0.053913	0.127451	0.099338	0.113333	0.072
Uinta	0.000000	0	0.345538	0.237718	0.197943	0.130682	0.283019	0.247154	0.224274	0.235
Washakie	14.836295	0	0.260163	0.195652	0.083333	0.076923	0.283186	0.139706	0.265306	0.194
Weston	0.000000	0	0.318841	0.241935	0.264706	0.266667	0.181818	0.290323	0.218750	0.285

2915 rows × 62 columns



Question 4

Create a separate dummy variable for each of the 48 states and the District of Columbia in the dataset (so you'll create 49 dummy variables in total, but dropping observations with missing values may reduce this number).

```
In [8]: dummies = pd.get_dummies(covid_data1['state'])
        covid_dummies = pd.concat([covid_data1, dummies], axis=1).reset_index()
```

```
In [9]: covid_dummies
```

```
Out[9]:
```

	county	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	...
0	Autauga	16.548864	1	0.333333	0.238095	0.208333	0.133333	0.375000	0.238095	0.260870	...
1	Baldwin	8.959118	1	0.268097	0.233503	0.167464	0.176991	0.298050	0.262467	0.193237	...
2	Barbour	6.609756	0	0.228571	0.250000	0.181818	0.111111	0.294118	0.571429	0.545455	...
3	Bibb	6.038192	1	0.244444	0.280000	0.181818	0.150000	0.466667	0.375000	0.190476	...
4	Blount	1.503713	1	0.304348	0.260870	0.352941	0.166667	0.347826	0.318182	0.529412	...
...
2910	Sweetwater	0.000000	0	0.349304	0.302658	0.227799	0.178484	0.247209	0.226943	0.254593	...

	county	deathspc	intersects_msa	cur_smoke_q1	cur_smoke_q2	cur_smoke_q3	cur_smoke_q4	bmi_obese_q1	bmi_obese_q2	bmi_obese_q3	...
2911	Teton	10.535728	0	0.263415	0.182018	0.092105	0.053913	0.127451	0.099338	0.113333	...
2912	Uinta	0.000000	0	0.345538	0.237718	0.197943	0.130682	0.283019	0.247154	0.224274	...
2913	Washakie	14.836295	0	0.260163	0.195652	0.083333	0.076923	0.283186	0.139706	0.265306	...
2914	Weston	0.000000	0	0.318841	0.241935	0.264706	0.266667	0.181818	0.290323	0.218750	...

2915 rows × 110 columns

Question 5

Split the sample into a training set (80%) and a test set (20%). Be sure to set a random seed so you can replicate your work.

```
In [10]: covid_dummies.drop(['county','state'], axis =1, inplace = True)

predictors = covid_dummies.columns.tolist()

print(predictors)
```

```
['deathspc', 'intersects_msa', 'cur_smoke_q1', 'cur_smoke_q2', 'cur_smoke_q3', 'cur_smoke_q4', 'bmi_obese_q1', 'bmi_obese_q2', 'bmi_obese_q3', 'bmi_obese_q4', 'exercise_any_q1', 'exercise_any_q2', 'exercise_any_q3', 'exercise_any_q4', 'brfss_mia', 'puninsured2010', 'reimb_penroll_adj10', 'mort_30day_hosp_z', 'adjmortmeas_amiall30day', 'adjmortmeas_chfall30day', 'med_prev_qual_z', 'primcare_revis_10', 'diab_hemotest_10', 'diab_eyeexam_10', 'diab_lipids_10', 'mammogram_10', 'cs00_seg_inc', 'cs00_seg_inc_pov25', 'cs00_seg_inc_aff75', 'cs_race_theil_2000', 'gini99', 'poor_share', 'inc_share_1perc', 'frac_middleclass', 'scap_ski90pcm', 'rel_tot', 'cs_frac_black', 'cs_frac_hisp', 'unemp_rate', 'cs_labforce', 'cs_elf_ind_man', 'cs_born_foreign', 'mig_inflow', 'mig_outflow', 'pop_density', 'frac_traveltime_lt15', 'hhinc00', 'median_house_value', 'ccd_exp_tot', 'score_r', 'cs_fam_wkidsinglemom', 'subcty_exp_pct', 'taxrate', 'tax_st_diff_top20', 'pm25', 'pm25_mia', 'summer_tmmx', 'summer_rmax', 'winter_tmmx', 'winter_rmax', 'bmcruderate', 'Alabama', 'Arizona', 'Arkansas', 'California', 'Colorado', 'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Idaho', 'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana', 'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota', 'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada', 'New Hampshire', 'New Mexico', 'New York', 'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon', 'Pennsylvania', 'Rhode Island', 'South Carolina', 'South Dakota', 'Tennessee', 'Texas', 'Utah', 'Vermont', 'Virginia', 'Washington', 'West Virginia', 'Wisconsin', 'Wyoming']
```

```
In [11]: #train, test = train_test_split(covid_dummies, test_size=0.2, random_state=10)

X= covid_dummies[predictors[1:]]
```

```

y= covid_dummies['deathspc']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=10)

X_train = pd.DataFrame(preprocessing.scale(X_train))
X_test = pd.DataFrame(preprocessing.scale(X_test))

```

Question 6

Using the training set, fit a model of COVID-19 deaths per capita ($y = \text{deathspc}$) as a function of the Opportunity Insights and PM COVID predictors listed in the spreadsheet, as well as state-level fixed effects (the state dummy variables) using OLS.

In [12]:

```

#X= covid_dummies[predictors[1:]]
#y= covid_dummies['deathspc']

#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=10)

#X_train = sm.add_constant(X_train)
#X_test = sm.add_constant(X_test)

X_train_norm = np.hstack((np.ones((X_train.shape[0], 1)), X_train))
model_norm = sm.OLS(y_train, X_train_norm).fit()
print(model_norm.summary())

y_train_pred = model_norm.predict(X_train_norm)
mse_train = mean_squared_error(y_train, y_train_pred)
print("MSE for the training set:", mse_train)

X_test_norm = np.hstack((np.ones((X_test.shape[0], 1)), X_test))
y_test_pred = model_norm.predict(X_test_norm)
mse_test = mean_squared_error(y_test, y_test_pred)
print("MSE for the test set:", mse_test)

```

OLS Regression Results

```

=====
Dep. Variable:      deathspc      R-squared:                0.432
Model:              OLS          Adj. R-squared:           0.405
Method:             Least Squares  F-statistic:             16.12
Date:               Wed, 22 Feb 2023  Prob (F-statistic):       1.63e-202
Time:               22:32:24      Log-Likelihood:          -11985.
No. Observations:   2332         AIC:                     2.418e+04

```

Df Residuals: 2226 BIC: 2.479e+04
Df Model: 105
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	23.3145	0.875	26.640	0.000	21.598	25.031
x1	1.8467	1.204	1.534	0.125	-0.514	4.207
x2	0.5730	1.454	0.394	0.694	-2.278	3.424
x3	-0.7181	1.272	-0.565	0.572	-3.212	1.776
x4	-0.8748	1.070	-0.818	0.414	-2.973	1.223
x5	0.4054	1.047	0.387	0.699	-1.648	2.459
x6	-0.5829	1.444	-0.404	0.686	-3.414	2.249
x7	0.7901	1.407	0.562	0.575	-1.969	3.550
x8	-2.8279	1.186	-2.384	0.017	-5.154	-0.502
x9	0.3917	1.125	0.348	0.728	-1.815	2.598
x10	-1.3011	2.146	-0.606	0.544	-5.510	2.908
x11	5.8906	2.536	2.322	0.020	0.917	10.864
x12	-1.6728	2.435	-0.687	0.492	-6.448	3.103
x13	-3.6157	2.576	-1.403	0.161	-8.668	1.436
x14	-3.2319	3.579	-0.903	0.367	-10.251	3.787
x15	-2.5734	2.486	-1.035	0.301	-7.448	2.301
x16	-2.4877	1.683	-1.478	0.140	-5.788	0.813
x17	0.1480	2.282	0.065	0.948	-4.326	4.622
x18	0.7174	1.644	0.436	0.663	-2.506	3.941
x19	0.1335	1.588	0.084	0.933	-2.980	3.247
x20	5.1777	4.574	1.132	0.258	-3.793	14.148
x21	-1.7907	1.474	-1.215	0.225	-4.681	1.100
x22	-6.4472	1.873	-3.442	0.001	-10.120	-2.774
x23	0.1356	1.878	0.072	0.942	-3.547	3.819
x24	-1.0490	2.004	-0.524	0.601	-4.979	2.881
x25	-1.1155	1.880	-0.593	0.553	-4.801	2.570
x26	39.1466	15.028	2.605	0.009	9.675	68.618
x27	-26.2382	7.442	-3.526	0.000	-40.832	-11.645
x28	-13.9143	8.482	-1.641	0.101	-30.547	2.718
x29	0.3628	1.305	0.278	0.781	-2.196	2.921
x30	-5.0277	2.443	-2.058	0.040	-9.819	-0.236
x31	1.5883	2.628	0.604	0.546	-3.565	6.741
x32	1.0183	1.779	0.572	0.567	-2.471	4.508
x33	-7.5686	2.148	-3.523	0.000	-11.781	-3.356
x34	-5.0455	1.897	-2.660	0.008	-8.766	-1.325
x35	2.7129	1.441	1.883	0.060	-0.112	5.538
x36	12.7262	2.463	5.167	0.000	7.896	17.556
x37	-1.3030	2.093	-0.622	0.534	-5.408	2.802

x38	-3.2877	1.543	-2.131	0.033	-6.313	-0.262
x39	-1.9666	2.038	-0.965	0.335	-5.963	2.029
x40	3.7595	1.461	2.573	0.010	0.894	6.625
x41	4.7141	2.011	2.344	0.019	0.770	8.658
x42	0.3743	2.120	0.177	0.860	-3.783	4.532
x43	-4.5086	1.967	-2.293	0.022	-8.365	-0.652
x44	18.3812	1.357	13.545	0.000	15.720	21.042
x45	-1.0386	1.732	-0.600	0.549	-4.434	2.357
x46	4.1132	2.734	1.504	0.133	-1.248	9.475
x47	0.8789	2.387	0.368	0.713	-3.803	5.561
x48	2.2306	1.237	1.803	0.071	-0.195	4.656
x49	1.9761	1.617	1.222	0.222	-1.194	5.147
x50	0.8865	2.530	0.350	0.726	-4.075	5.848
x51	-0.1006	1.229	-0.082	0.935	-2.511	2.310
x52	0.2795	1.452	0.192	0.847	-2.568	3.127
x53	0.5988	0.658	0.909	0.363	-0.692	1.890
x54	-1.8387	2.635	-0.698	0.485	-7.006	3.329
x55	-0.4794	1.102	-0.435	0.664	-2.640	1.681
x56	1.0306	3.063	0.336	0.737	-4.977	7.038
x57	-5.3779	3.422	-1.571	0.116	-12.089	1.333
x58	9.3939	4.773	1.968	0.049	0.034	18.754
x59	-1.3182	1.975	-0.668	0.504	-5.191	2.554
x60	0.8314	2.029	0.410	0.682	-3.148	4.810
x61	-2.8590	1.099	-2.602	0.009	-5.014	-0.704
x62	-2.5789	1.245	-2.071	0.038	-5.021	-0.137
x63	-3.8565	1.012	-3.811	0.000	-5.841	-1.872
x64	-7.0115	1.479	-4.740	0.000	-9.912	-4.111
x65	0.3000	1.239	0.242	0.809	-2.130	2.730
x66	4.6079	0.937	4.918	0.000	2.771	6.445
x67	0.8789	0.899	0.977	0.329	-0.885	2.643
x68	-4.5637	1.469	-3.106	0.002	-7.445	-1.683
x69	1.8338	1.439	1.274	0.203	-0.988	4.656
x70	-0.3684	1.401	-0.263	0.793	-3.116	2.380
x71	2.1191	1.142	1.855	0.064	-0.121	4.359
x72	7.4891	1.159	6.463	0.000	5.217	9.762
x73	2.2759	1.267	1.797	0.073	-0.208	4.760
x74	-0.6066	1.125	-0.539	0.590	-2.813	1.600
x75	1.3076	1.120	1.167	0.243	-0.889	3.504
x76	7.7762	1.295	6.003	0.000	5.236	10.317
x77	0.4000	1.035	0.387	0.699	-1.629	2.430
x78	0.2704	0.942	0.287	0.774	-1.576	2.117
x79	7.2678	0.980	7.414	0.000	5.345	9.190
x80	6.0944	1.159	5.257	0.000	3.821	8.368
x81	2.0004	1.414	1.414	0.157	-0.773	4.774

x82	-0.6219	1.250	-0.498	0.619	-3.073	1.829
x83	-0.4302	0.990	-0.434	0.664	-2.372	1.512
x84	0.8096	1.129	0.717	0.474	-1.405	3.024
x85	0.9028	1.092	0.827	0.408	-1.238	3.044
x86	-2.0615	1.187	-1.737	0.082	-4.389	0.265
x87	0.0174	0.933	0.019	0.985	-1.812	1.847
x88	-2.2463	1.288	-1.745	0.081	-4.771	0.279
x89	6.7316	1.214	5.544	0.000	4.350	9.113
x90	-3.7375	1.189	-3.144	0.002	-6.069	-1.406
x91	1.7203	1.267	1.358	0.175	-0.763	4.204
x92	1.8889	1.130	1.671	0.095	-0.327	4.105
x93	-1.3299	1.088	-1.223	0.222	-3.463	0.803
x94	-0.9965	1.258	-0.792	0.428	-3.464	1.471
x95	2.6481	1.097	2.413	0.016	0.496	4.800
x96	-0.5128	0.902	-0.569	0.570	-2.281	1.255
x97	-4.5864	1.100	-4.169	0.000	-6.743	-2.429
x98	0.9014	1.170	0.770	0.441	-1.393	3.196
x99	-1.8221	1.032	-1.765	0.078	-3.846	0.202
x100	-7.0452	1.878	-3.752	0.000	-10.728	-3.362
x101	-2.3114	1.454	-1.590	0.112	-5.162	0.540
x102	0.3696	0.978	0.378	0.705	-1.547	2.287
x103	-4.6794	1.127	-4.153	0.000	-6.889	-2.470
x104	0.6940	1.400	0.496	0.620	-2.051	3.439
x105	0.0501	1.025	0.049	0.961	-1.960	2.060
x106	0.9260	1.291	0.718	0.473	-1.605	3.457
x107	-1.1689	1.087	-1.075	0.282	-3.300	0.963

```

=====
Omnibus:                2270.523   Durbin-Watson:           2.091
Prob(Omnibus):          0.000   Jarque-Bera (JB):       166465.540
Skew:                   4.505   Prob(JB):               0.00
Kurtosis:               43.398   Cond. No.               1.64e+15
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 9.52e-27. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

MSE for the training set: 1704.9352722841477

MSE for the test set: 1476.592736354548

Part b

Why might you be concerned about overfitting in this context? Is there any evidence of overfitting? Briefly explain.

Above, we have actually fit 105 predictors on a sample of just 2332. In this context, overfitting could occur because the model seems too complex as it has a lot of parameters.

R-squared: R-squared is 0.432, which means that the predictors explain only 43.2% of variation in the dependent variable. Adjusted R-squared: This Adjusted R-squared value takes into account the number of independent variables in the model whose value is 0.405. This suggests that there may be some evidence of overfitting in the model. However, this may not be a right metric to observe overfitting. F-statistic: The F-statistic is 16.12 and the associated p-value is very small. This says that the model is statistically significant. AIC and BIC: These are information criteria that are used to compare different models. Lower values indicate better models. In this case, the AIC is $2.418e+04$ and the BIC is $2.479e+04$. These huge values suggests that the model may be overfitting, as it is penalized for including a large number of predictors. However, this may not be a right metric to observe overfitting. Thus from the above values from the regression results, i.e looking at Adjusted R squared, F-Statistic, AIC, BIC, we can say that the large number of predictors for such a small sample may be penalizing the model. This means that the model may be over stating the performance on training data set. However, the model may not generalize well to the new data.

However in the results we have observed, we see that the evidence in the form of MSE for training and test data set suggests that there is no overfitting. However, if we observe that there is a change in performance i.e decrease in performance on test set but on the other side the performance on training set increases, this can be taken as overfitting. To analyse this further, we can try to attempt to run cross validation either 5 or 10 fold, and detect for overfitting.

Question 7

Use the training set to estimate ridge regression and the lasso analogs to the OLS model in the previous question. For each, you should report a plot of the cross-validation estimates of the test error as a function of the value of the hyperparameter (λ) that indicates the tuned value of λ . Hint: to do so you should be sure standardize your predictors and tune the hyperparameter by:

```
In [13]: from sklearn.linear_model import RidgeCV, LassoCV
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
```

```
In [14]: #X= covid_dummies[predictors[1:]]
#X = pd.DataFrame(preprocessing.scale(X))
#y= covid_dummies['deathspc']

#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=25)
```

```
#X_train = pd.DataFrame(preprocessing.scale(X_train))
#X_test = pd.DataFrame(preprocessing.scale(X_test))
```

In [15]:

```
#Lasso
kf = KFold(n_splits=10, random_state = 25, shuffle=True)
tkf = kf.split(X_train, y_train)#change
lasso = Lasso()
alphacl = (np.logspace(-2, 2, 100))
alpha_lgrid = [{'alpha': alphacl}]
def vector_values(grid_search, trials):
    mean_vec = np.zeros(trials)
    std_vec = np.zeros(trials)
    i = 0
    final = grid_search.cv_results_

    for mean_score, std_score in zip(final["mean_test_score"], final["std_test_score"]):
        mean_vec[i] = -mean_score
        std_vec[i] = std_score
        i = i+1
    return mean_vec, std_vec

grid_search_lasso = GridSearchCV(lasso, alpha_lgrid, cv = tkf, scoring = 'neg_mean_squared_error')
grid_search_lasso.fit(X_train, y_train) #change
mean_vec, std_vec = vector_values(grid_search_lasso, 100)

results_cv = pd.DataFrame(
    {
        "alphas": alphacl,
        "MSE": mean_vec,
    }
)
min_mse = min(results_cv['MSE'])
results_cv.loc[results_cv['MSE']==min_mse]
plt.plot(
    alphacl,
    mean_vec,
    linewidth=2,
)

optimal_alpha = results_cv.loc[results_cv['MSE'] == min_mse]['alphas'].values[0]
```

```

print("MSE at optimal alpha: ", min_mse)


lasso_optimal = Lasso(alpha=optimal_alpha)
lasso_optimal.fit(X_train, y_train) #change

y_pred_lasso = lasso_optimal.predict(X_train) #change
mse_lasso = mean_squared_error(y_train, y_pred_lasso) #change
print("MSE for lasso with optimal alpha:", mse_lasso)

print("R square value for lasso with optimal alpha: " , lasso_optimal.score(X_train, y_train))

```

MSE at optimal alpha: 1927.757957321408
 MSE for lasso with optimal alpha: 1741.391623952136
 R square value for lasso with optimal alpha: 0.4198535762904305



In [16]:

```

# Ridge

#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=25)

#X_train = pd.DataFrame(preprocessing.scale(X_train))

kf = KFold(n_splits=10, random_state = 25, shuffle=True)
tkf = kf.split(X_train, y_train) #change
ridge = Ridge()

```

```

alphacr = (np.logspace(-2, 2, 100))
alpha_rgrid = [{'alpha': alphacr}]
grid_search_ridge = GridSearchCV(ridge, alpha_rgrid, cv = tkf, scoring = 'neg_mean_squared_error')
grid_search_ridge.fit(X_train, y_train) #change
mean_vec, std_vec = vector_values(grid_search_ridge, 100)
results_cv = pd.DataFrame(
    {
        "alphas": alphacr,
        "MSE": mean_vec,
    }
)
min_mse1 = min(results_cv['MSE'])
results_cv.loc[results_cv['MSE']==min_mse1]
plt.plot(
    alphacr,
    mean_vec,
    linewidth=2,
)

optimal_alpha1 = results_cv.loc[results_cv['MSE'] == min_mse1]['alphas'].values[0]
ridge_optimal = Ridge(alpha=optimal_alpha1)
ridge_optimal.fit(X_train, y_train) #change

print("MSE at optimal alpha: ", min_mse1)

y_pred_ridge = ridge_optimal.predict(X_train)
mse_ridge = mean_squared_error(y_train, y_pred_ridge)
print("MSE for Ridge with optimal alpha:", mse_ridge)


print("R square value for Ridge with optimal alpha:" , ridge_optimal.score(X_train, y_train))

```

```

MSE at optimal alpha: 1919.5891999810708
MSE for Ridge with optimal alpha: 1718.4122366607544
R square value for Ridge with optimal alpha: 0.42750918297462714

```



Question8

Using the ridge regression and the lasso models you trained based on the optimal values of λ you found in the previous question, calculate and report the training and test set prediction errors (MSE) for each model. Did ridge regression and/or the lasso improve your prediction over OLS? Which model performs the best? Briefly explain which model you would recommend to the CDC and why.

In [19]:

```
#Lasso

y_pred_train_lasso = lasso_optimal.predict(X_train) #change
mse_train_lasso = mean_squared_error(y_train, y_pred_train_lasso) #change
print("MSE for lasso with optimal alpha on train data:", mse_train_lasso)
print("R square value for lasso with optimal alpha on train data: " , lasso_optimal.score(X_train, y_train))

y_pred_test_lasso = lasso_optimal.predict(X_test) #change
mse_test_lasso = mean_squared_error(y_test, y_pred_test_lasso) #change
print("MSE for lasso with optimal alpha on test data:", mse_test_lasso)
print("R square value for lasso with optimal alpha on test data: " , lasso_optimal.score(X_test, y_test))
```

```
MSE for lasso with optimal alpha on train data: 1741.391623952136
R square value for lasso with optimal alpha on train data: 0.4198535762904305
MSE for lasso with optimal alpha on test data: 1441.6756819004447
R square value for lasso with optimal alpha on test data: 0.0988812619607149
```

In [20]:

```
# Ridge

y_pred_train_ridge = ridge_optimal.predict(X_train)
mse_train_ridge = mean_squared_error(y_train, y_pred_train_ridge)
print("MSE for Ridge with optimal alpha on training data:", mse_train_ridge)
print("R square value for Ridge with optimal alpha on training data:" , ridge_optimal.score(X_train, y_train))

y_pred_test_ridge = ridge_optimal.predict(X_test)
mse_test_ridge = mean_squared_error(y_test, y_pred_test_ridge)
print("MSE for Ridge with optimal alpha on test data:", mse_test_ridge)
print("R square value for Ridge with optimal alpha on test data:" , ridge_optimal.score(X_test, y_test))
```

```
MSE for Ridge with optimal alpha on training data: 1718.4122366607544
R square value for Ridge with optimal alpha on training data: 0.42750918297462714
MSE for Ridge with optimal alpha on test data: 1417.3555674076758
R square value for Ridge with optimal alpha on test data: 0.11408253861109552
```

MSE for lasso with optimal alpha on train data: 1741.391623952136

MSE for lasso with optimal alpha on test data: 1441.6756819004447

MSE for Ridge with optimal alpha on training data: 1718.4122366607544

MSE for Ridge with optimal alpha on test data: 1417.3555674076758

OLS - MSE for the training set: 1704.9352722841477

OLS - MSE for the test set: 1476.592736354548

Comparing the mean squared errors (MSE) obtained from the three models, we can see that, for both ridge regression and lasso improved the prediction over OLS. The MSE for both the ridge and lasso models on the test set is lower than the MSE for OLS on the same set, indicating better performance.

Based on the mean squared error (MSE) results provided, it appears that Ridge regression performs the best out of the three models on the test set. Hence I would recommend Ridge regression for this particular data set. Theoretically, the Ridge adds a penalty term to the OLS and minimizes the coefficients to close to zero resulting in a balance between the bias and variance. This will help with the problem of overfitting and also improves performance on any new data that we test on. Also due to the sensitivity of the data set i.e COVID death prediction, I would want to opt for a model

that performs better or you could say best of the options that are available. Thus I would choose the Ridge model based on an extensive analysis and tactical comparison of OLS, Lasso and Ridge based on the MSE values we have observed.

In []: